

Task 1 Machine learning on tabular mushrooms

The aim of this first task is to use machine learning techniques on a tabular dataset to determine if a mushroom is edible or poisonous. The dataset contains various attributes that describe the physical characteristics of mushrooms, such as shape, color, odor, and habitat. To ensure accurate classification, we have utilized various data preprocessing techniques, model selection, and evaluation metrics.

During our exploration of the dataset, we noticed missing values represented by "?" in the "stalk root" attribute. Instead of removing instances with missing values, which may result in data loss, we decided to treat the missing values as a separate category called "unknown." This allowed us to incorporate the unknown stalk root information into our classification models.

To convert categorical data into a suitable format for machine learning algorithms, we opted for one-hot encoding. This technique creates binary features for each category within the original feature, with 1 indicating the presence of the category and 0 indicating its absence. Although we initially considered using LabelEncoder, which assigns a numerical label to each category, one-hot encoding was a better choice because it avoids implying any ordinal relationship between categories that may not exist.

Once we preprocessed the data, we split it into a training set and a test set using the `train_test_split` function from the `sklearn` library. This function randomly divides the dataset into training and test subsets, ensuring that both sets have a similar distribution of instances. We allocated 80% of the data for training and 20% for testing.

We experimented with three different machine learning models: K-Nearest Neighbors (KNN), Logistic Regression, and Neural Network. Each model has unique characteristics and assumptions that can affect their performance on the given dataset.

K-Nearest Neighbors (KNN) is a powerful yet straightforward algorithm that classifies instances based on their proximity to the K nearest instances in the training set. The primary concern when using KNN is selecting an appropriate value for K, which affects the model's bias-variance trade-off. We chose $K=5$ as a reasonable starting point for our analysis.

Logistic Regression is a linear model for binary classification that estimates the probability of an instance belonging to a specific class. It assumes that the relationship between the features and the log odds of the target variable can be modeled using a linear function.

Logistic Regression inherently includes L2 regularization, which helps to prevent overfitting by penalizing large weights in the model.

Neural Networks are a class of models that can learn complex non-linear relationships between features and the target variable. They consist of interconnected layers of nodes, with each layer transforming the input data into a higher level of abstraction. The choice of the architecture, such as the number of layers and nodes, can significantly impact the model's performance. We used a relatively simple neural network architecture with two hidden layers, each containing 128 nodes, suitable for our dataset. After training each model on the preprocessed data, we evaluated their performance on the test set using accuracy as the primary evaluation metric. While accuracy is a widely used

metric, it may not be suitable for all classification problems, particularly when dealing with imbalanced datasets. Other evaluation metrics, such as precision, recall, F1-score, or the area under the Receiver Operating Characteristic (ROC) curve, may provide a more comprehensive understanding of the models' performance.

To further enhance the classification pipeline, we can consider the following steps and

Techniques:

1. Feature selection or dimensionality reduction: This involves using techniques such as Principal Component Analysis (PCA) or Recursive Feature Elimination (RFE) to reduce the number of features while still retaining the most relevant information.
2. Hyperparameter tuning: We can systematically search for the best hyperparameters by using techniques such as grid search or libraries like Optuna. This can lead to optimized model performance.
3. Cross-validation: By employing k-fold cross-validation, we can obtain a more accurate estimate of model performance and reduce the likelihood of overfitting.
4. Ensemble methods: Combining multiple models, such as bagging, boosting, or stacking, can improve overall performance and increase the robustness of predictions.
5. Model interpretability: We can use techniques such as SHapley Additive exPlanations (SHAP) or Local Interpretable Model-agnostic Explanations (LIME) to help explain the predictions of the models and identify the most important features.

In conclusion, our analysis demonstrates a comprehensive approach to machine learning on tabular data for the task of classifying mushrooms as edible or poisonous. By incorporating additional steps and considerations like feature selection or dimensionality reduction, hyperparameter tuning, cross-validation, ensemble methods, and model interpretability, we can enhance the classification pipeline and achieve a more robust and accurate solution. This approach can be applied to other classification tasks, providing valuable insights and predictions for various domains.

Task 2 Sentiment analysis

We started working on task 2 since the most recent lab that was out was lab3 which covered sentiment analysis. We started by implementing the BernoulliNB module by using the training data from train.json and creating our own test data by taking 20% from the train.json which gave us accuracy results with the accuracy being: Accuracy: 0.5981191222570533. This was before we realized that there is test.json.

Using the test.json and train.json, we then get a result of Accuracy: 0.7375105842506351. This is a significant increase, just by using the correct data. This happens because it allows the model to analyze more data to train itself to recognize the sentiment of words.

The next, and quite simple step, is to look at the data to see if there are any changes that can be made to help the model. For example, this would include the removal of “periods”, “commas”, “forward-slashes” and even unnecessary spaces. We would deem these characters unnecessary for the model because of the goal we have with it. In this task, we want to find and

count similar words within sentences to then set them with a defined sentiment. Removing these characters allows similar words to be counted together. As an example, before the removal, “good” and “good.” are considered two different words. This hinders the accuracy of the model quite less than expected.

Another thing we thought was important is to remove all capitalization so that words, as an example, “Good” and “good” become equal for the model.

After removing some unnecessary characters and removing capitalizations, the model then performed with Accuracy: 0.7383573243014394. Is it then even worth implementing?

I am unsure what the next step would be. Perhaps the idea is to continue looking for unnecessary characters to try and remove them. There must be a better way to improve the model. Perhaps another model may work better?

While testing our current model we found out that feeding the model correct data gave us higher results and removing punctuations and capitalizations adjusted it in a non-significant way.

We found out that by switching the model to multinomial model we got a far greater accuracy to the model. [Accuracy here](#)

The reason for multinomial model being better is because it takes into considers the frequency of the words compared to Bernoulli that look at if the word is present or absent

NEW, rewritten task 2. Not sure if we should have our original mistakes in, like using our training data also as the test data, as well as fitting both the test and training data

Looking over task two, we realized it seemed to be a very similar sentiment analysis task like in the third lab exercise. We took this into account and tried using the same methods as in the third lab. The two methods shown in the lab are using the models Bernoulli Naive Bayes and Multinomial Naive Bayes. This is done by taking each sentence and the positive, negative, or neutral attributes and connecting it using two separate lists, lining up the indexes. This is done to the training and testing data. The next step here is to vectorize the sentences list, meaning it turns the sentences into a numerical representation, so that the model can take the texts and sentiments as inputs. Initially, with Bernoulli Naive Bayes, without any text alteration, we got an accuracy score of 59.36%. This isn't great. There must be a way to get a higher accuracy score.

The first idea was to preprocess the training and testing data before predicting again with the test data. Removing some characters such as periods, commas, question marks, and unnecessary spaces, as well as making everything lowercase, yielded an accuracy score of, surprisingly, the same 59.36%.

We then thought that perhaps the other model might perform much better. We then tried the Multinomial Naive Bayes model with the same text preprocessing, getting an accuracy score of 61.39%. We were expecting better results.

The next thought was to continue looking at the third lab assignment. There it mentions TF-IDF, meaning Term Frequency - Inverse Document Frequency. TF measures how frequently a word is in the document. IDF measures how important a word is in comparison to the rest of the document. Implementing this only lowered the Multinomial model accuracy score slightly. It's all a little confusing. It seems that there must be some way to increase the accuracy score, but perhaps there just might not be enough data.

Task 3 Convolutional neural networks

In this task, the assignment is about training a convolutional neural network (CNN) as a binary classifier from the dataset that we have been provided with. This is a CIFAR-10 dataset that consists of 60000 images that are 32x32 colored images and will identify one of the following categories; airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck. Each of these has 6000 images that are divided into 50000 for the training model and 10000 for the testing model.

This is the last task we started working on since we have not yet covered the material in the curriculum until later. While researching for this assignment we found a couple of different pre-trained CNN models we could use for this project. Some of the choices were ResNet-50/101/152, Inception-v3/v4, MobileNetV1/V2/V3. A decision was made not to use these because of the computing power that would be unsatisfactory for some systems with the amount of layers that the models have. We landed upon the VGG16 model that is a deep CNN model that has 16 layers. It is a less complex model compared to the previous one we have mentioned, but it is more complex than a 3-5 layer CNN. We chose VGG16 at first because we ran into some errors while trying to create our own 3-5 layer CNN using a pre-trained CNN model that has strong performance in particular image classification.

As we implemented the model and chose our own pictures to test on the model we got an initial accuracy score of around 60.89% this wasn't just noticeable when looking at the accuracy. When feeding the model new images for it to classify it mis-categorized them in some instances, but was correct in other instances most of the time. Apparently, this might be a coincidence but worth mentioning that it was less accurate when we started mixing the files like using a .png and a .jpg. When we only used .png then it was accurate most of the time. The accuracy of the model is not only affected by the difference in the format of the input images but in the model's architecture, the training data, and the fine-tuning.

To improve the model's performance we increase the number of training epochs we are using from 10 which gives the initial accuracy to 15. This gave us an accuracy of 67.13 but we have to be aware that adding more epochs makes it so that it trains longer and requires more computing power and we have to be aware not to overfit our model either. So finding the appropriate amount is important. We could also adjust the hyperparameter tuning of the learning rate or batch size. We also could fine tune more layers of the model to increase its score

When using the model to classify new images that we have taken or decided to use, it is important to preprocess them in the same way as we did with the training data. This would include resizing the images, normalizing their pixel values, and ensuring that they are in the same format. Doing this leads to the model doing their job great of generalizing its learning to new data and maintaining a consistent performance.

In conclusion, this assignment demonstrates the usage and potential of a deep convolutional neural network, such as VGG16, for image classification tasks. Although our initial accuracy was not very high, adjusting different parameters led to improving the model.